

Real-Time Search

*Notes due to – Brian Roberts, grad. student at CS Dept., WPI
MAS Course*

Introduction

To search, according to Merriam-Webster's online dictionary, means “to look into or over carefully or thoroughly in an effort to find or discover something: as to examine in seeking something; to make painstaking investigation or examination.” These definitions accurately describe the term, as used in the multi-agents systems field. An agent, in an initial state or position, must carefully investigate possible options, in order to find a path to a final goal or solution.

Although there are many different algorithms available, these can be categorized into two main classes: off-line and real-time. An off-line search finds a complete solution path before attempting any steps towards the goal. In contrast, a real-time search computes just enough, so as to determine a reasonable next step. After taking the previous 'most reasonable' move, a real-time search algorithm performs further computation to determine the next 'most reasonable' step...and so on, repeating until a final state or goal is reached. Real-time searches do not guarantee to find an optimal solution, but serve only to interleave planning steps and actually executing them.

Put simply, these searches do what is known as 'path-finding'. This is used when an agent needs to get from one state to a (possibly) distant state through a sequence of operators or moves. A path-finding problem, MAS and other areas, consists of the following components:

- ? N : a set of nodes, each representing a state,
- ? L : set of directed links, each representing an operator available to a problem-solving agent.
- ? s : a unique node, representing the start node
- ? G : set of nodes G , each representing a final state ($G \subseteq N$)

Every directed link, l , has an associated weight or distance, which represents the cost to the agent of applying the operator. Nodes which can be reached directed from a node i , are called the *neighbors* of node i .

The basis of the algorithms mentioned above, and later in this text, is asynchronous dynamic programming. This process relies on the *principle of optimality*: a path P is optimal if, and only if all sub-paths along the path are optimal. That is, for every intermediate path between the starting and ending node, the path must be the shortest possible available.

For convenience, $h^*(i)$ is used to denote the distance from node i to the nearest goal node.

Asynchronous dynamic programming repeatedly computes successive approximations of $h^*(i)$, denoted by $h(i)$, at each node in the graph. For true asynchronous dynamic programming, every node must contain an asynchronous process, independently updating its $h(i)$ approximation function. Each node i follows the following algorithm:

Compute for each neighbor node j :

$$f^*(j) = k(i, j) + h(j)$$

Where:

$k(i, j)$: distance from node i to j

$h(i)$: estimated shortest distance from node i to goal

Update $h(i) = \min_j f(j)$

This process is repeated until the approximations $h(i)$ convert to the actual values of $h^*(i)$. It can be proved that $h(i)$ will converge for every node i , provided that $k(i, j)$ never over-estimates the cost.

Algorithms

A*

No paper would be complete without mentioning A* (pronounced: “A star”), perhaps the most famous heuristic off-line searching algorithm of all-time. Several real-time algorithms have been based off of A*, including the Learning Real-Time A* and Real-time A* algorithms.

The basic concept of the A* algorithm is a best-first search—the most probable paths are explored first, searching outward from the starting node until it reaches the goal node. The best path is determined by choosing the option with the lowest cost, where cost is measured by the function: $f(n) = g(n) + h'(n)$. The function $g(n)$ is the actual cost of the path so far, while $h'(n)$ is a heuristic function of the estimated cost of the path from the current state to the final goal.

This algorithm will always find an optimal path to a goal, provided the heuristic function is 'admissible'—it does not over-estimate the cost from the current node to the goal. It is worthy to note, that if the heuristic function too often under-estimates the costs, then A* will have to search too many probable paths, and may run out of time before any solution is found. However, if the heuristic function too often over-estimates the costs, a solution can be found very quickly, but will likely be sub-optimal. The admissibility of $h'(n)$ is error-tolerant—should $h'(n)$ occasionally over-estimates h of the path by cost C , then A* will only occasionally find an answer that costs more than cost C more than the optimal cost.

Learning Real-Time A*

The algorithm **Learning Real-Time A*** (LRTA*) is conceptionally almost the same as A*; LRTA* can be viewed as the real-time version of A*. The traversal through the nodes is as follows:

for each neighbor j of current node i

compute $f(j) = k(i, j) + h(j)$

store the new $h(i) = \min_j f(i)$

move to the node j with $\min f(j)$

Ties for minimal $f(j)$ values are broken randomly. LTRA* 'learns' the optimal path via repeated trails. As long as initial $h(i)$ values are not over-stated, this algorithm is guaranteed to find a solution. Over repeated trials, LTRA* will find the most optimal solution. It should be noted, that using the proper heuristic estimation function will speed up finding the optimal solution. A simple example of the difference good heuristics can make can be observed in the 8-puzzle game. The 8-puzzle consists of 8 numbered, movable tiles, in a 3x3 frame. One of the cells of the frame is empty, which makes it possible to move the tiles. The object of the puzzle is to transform one configuration of the tiles into another. Two well-known heuristic estimation functions for this problem are: the number of misplaced tiles, and the sum total of the Manhattan distances of misplaced tiles.

The first function is easily understood: if x tiles are incorrectly placed, then the estimation function $h(i)$ would return x . The second formula is only slightly harder to grasp. A Manhattan distance measurement is equated by summing the horizontal and vertical distances of each tile from its goal position. Thus, if the tiles are currently in the following position:

1	5	4
	3	8
2	6	7

Current state

And the goal is to arrange the tile in this manner:

1	2	3
4	5	6
7	8	

Goal State

Then, since the '1' tile is in the correct position, its Manhattan distance is 0. The '2' is shifted 1 position horizontally and 2 positions vertically from its goal position, adding up to a distance of 3.

Tile	Distance		
	Horizontal	Vertical	Manhattan
1	0	0	0
2	1	2	3
3	1	1	2
4	2	1	3
5	0	1	1
6	1	1	2
7	2	0	2
8	1	1	2
Total	8	7	15

By counting the number of non-zero manhattan distances, we arrive at the number of misplaced tiles in this example: 8.

A good heuristic function should estimate the cost of moving from the current state to the goal state, as accurately as possible, but without over estimating the distance. Since every tile must be moved individually, the Manhattan displacement of a tile will never exceed the number of moves required to move it into its correct position. The total cost of moving all the tiles into their correct positions can never exceed the sum total Manhattan distance. The misplaced tiles metric can be calculated by counting the number of non-zero Manhattan numbers, and therefore will never be greater than the Manhattan distance. From the above, we derive the following relationships:

$$misplaced_tiles(i) \leq manhattan(i) \leq h(i)$$

This clearly shows that the Manhattan distance is a better estimation to use as a heuristic function.

Real-Time A*

Real-Time A* (RTA*), is exactly the same as LRTA*, with one exception. When the new value of $h(i)$ is stored, it is updated with the second least value. So, the algorithm for RTA* reads:

for each neighbor j of current node i

compute $f(j) = k(i,j) + h(j)$

store the new $h(i) = \text{secondmin}_j f(j)$

move to the node j with $\min f(j)$

The key difference between RTA* and LRTA* (updating $h(i)$ with $\text{secondmin}_j f(j)$ versus $\min_j f(j)$) allows an agent using the RTA* to make locally-optimal decisions on tree structures. That is, the path toward the goal chosen by RTA* has the minimum estimated cost based on the already-obtained information. Although this allows RTA* to overestimate heuristic costs, RTA* learn more efficiently on a single trial, and will usually arrive at the goal quicker. If a problem is to be solved repeatedly with the same goal state but different initial states, LRTA* will improve its performance over time, and eventually converge on the optimal solutions to along all paths.

Moving Target Search

Moving Target Search (MTS) allows one to relax the assumption that the goal location remains stationary. This allows the target to move while the search is in progress. The goal is reached when both the solver and the target are standing on the same node in the graph.

Several variables are used in this search algorithm: x_i is the current position of the problem solver, x_j is neighbor positions of problem solver, y_i is the current position of target, and y_j is the neighbor positions of target.

The algorithm has two parts, one part to be executed when the problem solver moves, the other to be executed when the target moves.

When problem solver moves

For each neighbor x_j

calculate $h(x_j, y_i)$

Update $h(x_i, y_i) = \max(h(x_i, y_i), \min_{x'j} h(x'j, y_i) + 1)$

Move to x_j where $\min h(x_j, y_j)$ and update x_i

When target moves:

Calculate $h(x_i, y_j)$ for target's new y_j

$h(x_i, y_i) = \max(h(x_i, y_i), h(x_i, y_j) - 1)$

Update position $y_i = y_j$

Real-Time Bidirectional Search

Real-time Bidirectional Search (RTBS) involves two agents simultaneously searching in the same environment. There are two types of RTBS: centralized and decoupled. In centralized RTBS, agents coordinate their actions for mutual benefit. In decoupled RTBS, each agent chooses its actions

independently. One agent maybe even works against the other agent, albeit unintentionally. The algorithm is fairly simple, because it must use another, previous algorithm for the actual searching.

Repeated execute following:

Control step: select forward or backwards move

Forward move: Problem solver from initial state moves towards second, 'target' agent.

Backwards move: Solver from goal state moves towards first agent.

In **Centralized RTBS**, one algorithm to use is LRTA*/B, a bidirectional version of LRTA*. The algorithm is described as follows (adapted from [4], with modifications to allows costs > 1):

LRTA*/B

Control Strategy:

Calculate $h(x', y)$ for each neighbor x' of x .

Calculate $h(x, y')$ for each neighbor y' of y .

Update the value of $h(x, y)$ as follows:

$$h(x, y) \leftarrow \min_{x', y'} \left\{ \begin{array}{l} h(x', y) + k(x', y) \\ h(x, y') + k(x, y') \end{array} \right\}$$

If $\min_{x'} h(x', y) < \min_{y'} h(x, y')$, select a forward move.

If $\min_{x'} h(x', y) > \min_{y'} h(x, y')$, select a backward move.

Otherwise, select a forward or backward move randomly.

Forward move:

Move the forward problem solver to the neighbor x' with the minimum $h(x', y)$, i.e., assign the value of x' to x . (Ties are broken randomly.)

Backward move:

Move the backward problem solver to the neighbor y' with the minimum $h(x, y')$, i.e., assign the value of y' to y . (Ties are broken randomly.)

Similarly, the RTA* algorithm may be adopted for bidirectional search, resulting in the RTA*/B algorithm.

This strategy calculates all possible forward and backward moves, and selects the best choice. Both solvers are completely controlled by the centralized decision-maker. The control strategy of LRTA*/B clearly shows the coordination of the two agents: the solver whose $h(x, y)$ distance estimation for reaching the other agent is smaller is allowed to move. Since the move with the smaller cost total is always taken, LRTA*/B will take the shortest path which can be found with LTRA*.

In **Decoupled RTBS**, agents do not coordinate their moves, but independently seek to meet with the other agent. The only search algorithm which can handle this independence is the MTS algorithm. Ishida describes a bidirectional MTS, called MTS/B in [4]. The bidirectional MTS/B differs from the unidirectional MTS mainly in the fact that both agents are actively seeking to rendezvous with the other agent.

In the decoupled control strategy, the algorithm “employs the minimal control necessary to ensure the termination of the algorithm” [1]. This results in both problem solvers independently making all decisions. The algorithm described in both [1] and [4] simply select forward and backward moves alternately.

One surprising constraint of MTS/B is that the problem solver must move faster than the target. This assumption was originally introduced in the description of MTS, when the target was trying to avoid the problem solver. This assumption must remain in order to prevent the agents from infinitely looping. Without this assumption, infinite looping could occur when agents are on opposite sides of an obstacle. If each agent attempts to move around the obstacle, but end up either continually chasing each other around the obstacle, or shuffling back and forth from side to side on opposite ends of the barrier.

In addition to classifying RTBS algorithms into decoupled- and centralized-control, each algorithm can be further sub-divided along the lines of how much information is shared between agents. Agents can either share heuristic information or maintain and update their own $h(x, y)$ values. These classifications are called: **shared heuristic information** and **distributed heuristic information** schemes. MTS/B is a distributed heuristic decoupled-control RTBS algorithm, while LRTA*/B is a shared heuristic centralized-control RTBS algorithm.

The performance of both decoupled- and centralized control RTBS algorithms has been compared [1 and 4] against each other. When heuristic functions can be relied on to supply accurate values, the decoupled algorithm performed better. While the number of moves needed by the agents to rendezvous were approximately the same, the decoupled RTBS must only expand half as many states as the centralized version. This means that decisions can be made independently by the two problem solvers without losing efficiency, while halving the planning time needed. In uncertain situations, when the heuristic function cannot be relied upon to supply accurate values: the centralized strategy works best. The number of moves required by the centralized RTBS algorithm was approximately one-half that of its decoupled counterpart.

The performance of both RTBS algorithms was also compared to Real-time Unidirectional Search (RTUS). When the heuristic function is accurate, the difference in number of moves required in decoupled-RTBS, centralized-RTBS and RTUS is negligible. This is because, in clear situations, various search techniques can be used to obtain near optimal solutions. In uncertain situations, clear differences were observed, but can be generally applied. The comparison of RTBS to RTUS is highly reliant on the problem type. For n-puzzles, RTBS is clearly a better choice, reducing the number of moves over RTUS by a factor of 2 for 15-puzzles, and by 6 for 24-puzzles. However, in randomly generated mazes, RTUS becomes increasingly better as the obstacle-ratio increases.

Summary

Modern search algorithms are based on asynchronous dynamic programming and can be categorized into two categories: off-line and real-time. Real-time is used when movement steps must be interleaved with planning steps.

Many different algorithms for searching have been proposed, the most often used algorithms include: A*, Learning Real-Time A* (LRTA*), Real-Time A* (RTA*), and Moving Target Search (MTS). A* is an often used off-line algorithm. LRTA* is a real-time version of A*. RTA* differs from LRTA* only in updating its heuristic values. MTS allows the problem solver to seek a target that moves.

Bidirectional search is when 2 agents want to rendezvous. Real-Time Bidirectional Search (RTBS) is categorized into two different levels coordination: decoupled and centralized. The information sharing can also be classified into: distributed or shared. LTRA*/B is centralized and shared, while MTS/B is decoupled and distributed.

Bidirectional search works similarly to unidirectional when heuristic functions are accurate, but often differ when distances are uncertain. Bidirectional works better for n-puzzle problems, unidirectional for mazes with a high (>20%) obstacle-density.

References

Real-time bidirectional search: Coordinated problem solving in uncertain situations. T. Ishida, IEEE Trans. on PAMI, Vol.18, No.6, 1996, p.617-627.

Moving-Target Search: A Real-Time Search for Changing Goals. T. Ishida. and R. Korf. IEEE Trans. on PAMI, Vol.17, No.6, 1995, p.609-619.

The Trailblazer Search: A New Method for Searching and Capturing Moving Targets. T. Sasaki, et al. AAAI-94, pp. 1,347-1,352. 1994

Two is not Always Better than One: Experiences in Real-Time Bidirectional Search. T. Ishida. Int'l Conf. Multi-Agent Systems (ICMAS-95, pp. 185-192. 1995

Multi-Agent Systems: A Modern Approach to Distributed Artificial Intelligence. G. Weiss (ed). MIT Press, 2001. Chapter 4, pp. 165-199.